

NLP - Assignment 1

Nonlinear Regression with Deep Learning

Dhivya Chandrasekaran
Department of Computer Science)
Student ID : 1111017
dchandra@lakeheadu.ca

Abstract—In this assignment we intend to build a Non-linear Regression model using 1 Dimensional Convolutional Neural Network.

I. INTRODUCTION

Non linear regression is a form of regression analysis in which the decision boundary is not a straight line and rather fits the data in a more flexible curve. The aim of this assignment is to build a custom made 1 dimensional convolutional neural network(CNN) using PyTorch. PyTorch is an open source Python library widely used for Natural Language Processing(NLP) tasks. The article provides a brief overview of convolutional neural networks in section II and the experimental analysis of the model is carried out in section III and section IV provides an outline of the results.

II. BACKGROUND

The *one-dimensional convolution* is defined as an operation between a vector of weights and a vector of sequential inputs where the weights act as filters for the data. The idea behind convolution is the element-wise multiplication of these vectors to form another vector. Convolution layers are used to capture primitive features in a data. [1] *Maxpooling* is a non-linear subsampling technique that returns only the maximum value in a given set of input. Maxpooling layers are used in convolutional neural networks to reduce the dimensionality of the input data. *Convolutional neural networks (CNN)* are a class of deep learning models which used the operation of convolution in atleast one of its layers. Dr. Yann Lecun invented the LeNet which is the first deep learning model. Later various CNNs models were proposed of which the most prominent are

AlexNet, VGGNet, GoogleNet and ResNet. CNNs are widely used in Image Processing and have become the basis of modern Computer Vision [2]. We calculate the *R² score* for the model proposed in this article. R squared error can be defined as the measure of the closeness to the prediction values to the fitted regression line. The score varies between 0 and 1. 1 indicates that the predicted data points are closer to the expected values. We calculate the *L1 Loss* in our model which measure the absolute value difference between the desired output and the predictions.

III. EXPERIMENTAL ANALYSIS AND COMPARISON

The entire code for the assignment is provided in Appendix A

A. Dataset:

For this assignment the *California housing dataset* is used. This dataset contains the details from 1990 California census. The dataset contains 20,640 instances and 10 features. The 10 different features and the description are given in Table I. Missing data in the dataset is handled by elimination since only 207 instances have missing values.

B. Model building:

The code snippet for the proposed model is shown in A. The proposed model has three convolution layer followed by one max pooling layer. The output from the maxpooling layer is fed to a flatten layer in order to be to a linear layer. The linear layer output is passed to the output layer which provides the predictions. The proposed model also performs Batch Normalisation to handle vanishing gradients




Fig. 1. Line graph depicting features of the dataset

Features	Description
longitude	A measure of how far west a house is; a higher value is farther west
latitude	A measure of how far north a house is; a higher value is farther north
housingMedianAge	Median age of a house within a block; a lower number is a newer building
totalRooms	Total number of rooms within a block
totalBedrooms	Total number of bedrooms within a block
population	Total number of people residing within a block
households	Total number of households, a group of people residing within a home unit, for a block
medianIncome	Median income for households within a block of houses (measured in tens of thousands of US Dollars)
medianHouseValue	Median house value for households within a block (measured in US Dollars)
oceanProximity	Location of the house w.r.t ocean/sea

TABLE I
LIST OF FEATURES IN THE DATASET

and exploding gradients. The input data is passed through BatchNorm1d layer before passing through the first convolution layer. The model was trained using three different optimizers SGD, ASGD and Adam optimizer. Out of the three, Adam optimizer yielded the best results hence it was used in the final model. The model was evaluated with three different activation functions in the convolution layers of the model such as relu, sigmoid and tanh. Sigmoid and

TanH caused errors due to vanishing gradients hence ReLU activation function is used in the model. In order to take into consideration all the features available the kernel size is maintained at 1.

No of Epochs	Learning Rate	Training R2Score	Testing R2Score	Inference Time in ms
50	0.01	0.719428	0.669038	182
100	0.01	0.732206	0.721175	141
300	0.01	0.758559	0.744018	139
500	0.01	0.777081	0.727654	229
1000	0.01	0.779681	0.707925	145

TABLE II
PERFORMANCE COMPARISON AT DIFFERENT NO OF EPOCHS
WITH FIXED LEARNING RATE

No of Epochs	Learning Rate	Training R2Score	Testing R2Score	Inference Time in ms
300	0.1	0.684762	0.689833	147
300	0.01	0.758559	0.744018	139
300	0.001	0.753317	0.732106	147

TABLE III
PERFORMANCE COMPARISON AT DIFFERENT NO OF EPOCHS
WITH FIXED LEARNING RATE

Table II provides the comparison of performance of the model when implemented for different number of epochs at a constant learning rate. Table III provides comparison of performance of the model

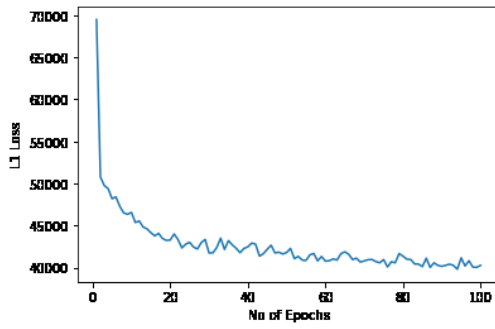


Fig. 2. Loss Curve

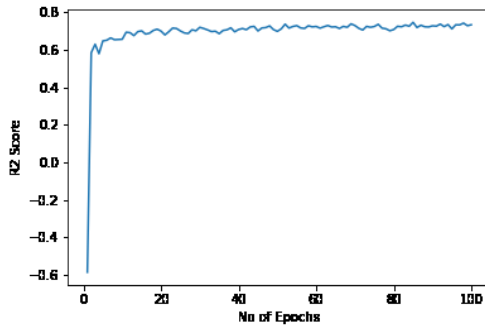


Fig. 3. R2 Score Curve

with different learning rates executed with the most effective number of epochs from Table II. Based on the tabulated results the optimal learning rate for the model is 0.01 and the best results are produced when repeated for 300 epochs.

IV. RESULTS

The data is mounted using the code shown in A and the code to display the first ten records of the dataset is shown in C. The output line graph is shown in Fig.1. The optimal solution for the model is obtained using Adam Optimizer, ReLU activation function, learning rate 0.01, repeated for 300 epochs. The training r2 score is . The testing r2 score is 0.744018 and the L1 Loss is 38706.48 . The total number of trainable parameters in the model are 74,961. The best inference time achieved is 139ms. The loss curve and the r2 score curve are shown in Fig 2 and Fig 3 respectively.

REFERENCES

- [1] Hu, Baotian, Zhengdong Lu, Hang Li, and Qingcai Chen. "Convolutional neural network architectures for matching natural language sentences." In Advances in neural information processing systems, pp. 2042-2050. 2014.

- [2] Gatys, Leon A., Alexander S. Ecker, and Matthias Bethge. "Image style transfer using convolutional neural networks." In Proceedings of the IEEE conference on computer vision and pattern recognition, pp. 2414-2423. 2016.

APPENDIX

A. Mounting data to Google Colab

```
1 #Mount the google drive and provide authentication
2 from google.colab import drive
3 drive.mount('/content/drive')
```

B. Importing necessary libraries for data processing

```
1 from sklearn import linear_model
2 from sklearn.model_selection import train_test_split
3 from sklearn.metrics import mean_squared_error, r2_score
4 import pandas as pd
5 import numpy as np
6 import matplotlib.pyplot as plt
```

C. Importing necessary libraries for data processing

```
1
2 dataset = pd.read_csv('/content/drive/My Drive/housing.csv')
3 print("The first ten rows of the dataset are:")
```

D. Read dataset and print the first 10 rows

```
1 dataset = pd.read_csv('/content/drive/My Drive/housing.csv')
2 print("The first ten rows of the dataset are:")
3 dataset.head(10)
```

E. Plot each feature of the dataset on separate subplots

```
1 axes = dataset.plot.line(subplots=True, figsize=(20, 10))
2 plt.xlabel("Number of Instances")
```

F. Data cleaning by removing null values

```
1 dataset = dataset.dropna()
```

G. Split data into features and predictions and split the data into training and test in the ratio 70:30

```
1 Y = dataset['median_house_value']
2 X = dataset.loc[:, 'longitude': 'median_income']
3
4 # Split the data into training and test sets with 70% training samples and 30% test samples
```

```

5 x_train, x_test, y_train, y_test =
    train_test_split(X, Y, test_size=0.3,
        random_state=2003)

```

H. Convert the data to numpy array to be compatible with PyTorch

```

1 x_train_np = x_train.to_numpy()
2 y_train_np = y_train.to_numpy()
3 x_test_np = x_test.to_numpy()
4 y_test_np = y_test.to_numpy()

```

I. Import necessary libraries for model execution using Pytorch

```

1 import time
2 import torch
3 from torch.nn import Conv1d
4 from torch.nn import MaxPool1d
5 from torch.nn import Flatten
6 from torch.nn import Linear
7 from torch.nn import BatchNorm1d
8 from torch.nn.functional import relu #
    activation function
9 from torch.utils.data import DataLoader,
    TensorDataset
10 from torch.optim import SGD
11 from torch.optim import Adam
12 from torch.optim import ASGD
13 from torch.optim import Adamax
14 from torch.nn import L1Loss
15
16 !pip install pytorch-ignite
17 from ignite.contrib.metrics.regression.
    r2_score import R2Score

```

J. Initialize and define the CNN model by overwriting the CnnRegressor Class

```

1 #Initialize the superclass
2 super(CnnRegressor, self).__init__()
3 self.batch_size = batch_size
4 self.inputs = inputs
5 self.outputs = outputs
6
7 self.input_layer = Conv1d(inputs,
    batch_size, 1)
8 self.batch_norm = BatchNorm1d(8)
9 #First Convolution Layer: Filters 128
10 self.conv_layer = Conv1d(batch_size, 128, 1)
11 #Second Convolution Layer: Filters 128
12 self.conv_layer_2 = Conv1d(128, 128, 1)
13 #Third Convolution Layer: Filters 256
14 self.conv_layer_3 = Conv1d(128, 256, 1)
15 #Max pool layer
16 self.max_pooling_layer = MaxPool1d(1)
17 #Flatten layer
18 self.flatten_layer = Flatten()
19 #Linear layer
20 self.linear_layer = Linear(256, 64)
21 #Output layer
22 self.output_layer = Linear(64, outputs)
23

```

```

24 def feed(self, input):
25
26     input = input.reshape((self.batch_size,
        self.inputs, 1))
27     output = self.batch_norm(input)
28     output = relu(self.input_layer(output))
29     #Three convolution layers
30     output = relu(self.conv_layer(output))
31     output = relu(self.conv_layer_2(output))
32     output = relu(self.conv_layer_3(output))
33     #max pooling layer
34     output = self.max_pooling_layer(output)
35     #Flatten layer
36     output = self.flatten_layer(output)
37     #Linear Layer
38     output = self.linear_layer(output)
39     #Output Layer
40     output = self.output_layer(output)
41
42     return output #return the output

```

K. Build the model

```

1
2 # Build the defined model with batch size 64
3
4 batch_size = 64
5 model = CnnRegressor(batch_size, X.shape[1], 1)
6 model.cuda() #to use GPU

```

L. Define model_loss function to calculate the performance metrics

```

1 # Function to calculate the average L1 Loss
    and R2 Score for every epoch.
2
3 def model_loss (model, dataset, train = False,
    optimizer = None):
4     performance = L1Loss()
5     score_metric = R2Score()
6
7     avg_loss = 0
8     avg_score = 0
9     count = 0
10
11     for input , output in iter(dataset):
12         predictions = model.feed(input)
13
14         loss = performance(predictions, output)
15
16         score_metric.update([predictions, output])
17         score = score_metric.compute()
18
19         if(train):
20             optimizer.zero_grad()
21
22             loss.backward()
23
24             optimizer.step()
25
26         avg_loss += loss.item()
27         avg_score += score
28         count += 1

```

```

29
30 return avg_loss / count, avg_score/count

```

M. Train the model

```

1 epochs = 300
2
3 optimizer = Adam(model.parameters(),lr=0.001)
4
5 torch.cuda.synchronize()
6 tr_st_time = int(round(time.time()*1000))
7
8 inputs = torch.from_numpy(x_train_np).cuda().
    float()
9 outputs = torch.from_numpy(y_train_np.reshape(
    y_train.shape[0],1)).cuda().float()
10
11 r2_score_plot = []
12 loss_plot = []
13
14 tensor = TensorDataset(inputs,outputs)
15 loader = DataLoader(tensor,batch_size, shuffle
    =True,drop_last=True)
16
17 for epoch in range(epochs):
18     avg_loss, avg_r2_score = model_loss(model,
        loader,train=True,optimizer=optimizer)
19     r2_score_plot.append(avg_r2_score)
20     loss_plot.append(avg_loss)
21     print ("Epoch" +str(epoch+1)+" :\n\tLoss =" +
        str(avg_loss)+"\n\tR2_Score =" +
        str(avg_r2_score))
22
23
24 torch.cuda.synchronize()
25 tr_time_elapsed = int(round(time.time()*1000))
    - tr_st_time

```

N. Save the model

```

1 torch.save(model.state_dict(), '111101
    _ldconv_reg.pt')

```

O. Plot the results

```

1
2 x_label = np.arange(1, (epochs+1))
3 plt.figure()
4 plt.xlabel("No of Epochs")
5 plt.ylabel("L1 Loss")
6 plt.plot(x_label,loss_plot)
7 plt.figure()
8 plt.xlabel("No of Epochs")
9 plt.ylabel("R2 Score")
10 plt.plot(x_label,r2_score_plot)
11 plt.show()

```

P. Test the model

```

1 inputs = torch.from_numpy(x_test_np).cuda().
    float()
2 outputs = torch.from_numpy(y_test_np.reshape(
    y_test_np.shape[0],1)).cuda().float()
3
4 tensor = TensorDataset(inputs,outputs)

```

```

5 loader = DataLoader(tensor,batch_size, shuffle
    =True,drop_last=True)
6
7 torch.cuda.synchronize()
8 ts_start_time = int(round(time.time()*1000))
9 avg_loss, avg_r2_score = model_loss(model,
    loader)
10
11 torch.cuda.synchronize()
12 ts_time_elapsed = int(round(time.time()*1000))
    - ts_start_time

```

Q. Print the results

```

1 print("Training time elapsed: {} ms".format(
    tr_time_elapsed))
2 print ("Test time elapsed: {} ms".format(
    ts_time_elapsed))
3
4 print("The model's L1 loss is: %.2f " %
    avg_loss)
5 print("The model's R^2 score is : %.6f" %
    avg_r2_score)
6 model_parameters = filter(lambda p: p.
    requires_grad, model.parameters())
7 params = sum([np.prod(p.size()) for p in
    model_parameters])
8 print("The number of trainable parameters =",
    params)

```